

SKOLKOVO INSTITUTE OF SCIENCE AND
TECHNOLOGY

MA06018 (TERM 3, 2018-2019)

Online ARIMA

Submitted By :

Bespalov Iaroslav

Vyacheslav Kozitsin

Iurii Katser

Edgar Makarov

Contents

Introduction	2
Related Work	3
Dataset Description	4
ML Methods and Experiments	6
Conclusion	9

Introduction

The ubiquity of time series is a fact of modern-day life: from stock market data to social media information; from search-engine metadata to webpage analytics, modern-day data exists as a continuous flow of information indexed by timestamps.

ARIMA models are the main time-series long-term forecasting models in many practical applications. In time series trends or dependencies usually change. This requires frequent refitting of the model from scratch or performing additional training. However, fitting ARIMA models to large-scale data (e.g. hourly time series for 10 years) is time consuming. In this case, the online learning approach helps.

For a real-time application, online modeling is desirable — a self-generating and self-updating forecasting method is desirable for producing a timely forecast. The ARIMA method in time series is an effective means to address a forecast; it has a solid foundation in classical probability theory and mathematical statistics [1].

We propose a comparison of different methods of seasonal ARIMA online implementation for two separated datasets. Therefore, the objective of this study was to investigate time complexity and performance of different online learning approaches:

- complete refitting on data of steps $[1, t]$ each step t ;
- rolling-window refitting on $[t-h, t]$;
- additional fitting (warm start) for every new observation.

Related Work

In literature, very few study has seriously investigated scalable algorithms for ARIMA models, although some simple methods were attempted. For example, the iterated least-squares approach was developed to consistently estimate autoregressive parameters [2]. Least squares and gradient algorithms are presented through estimating residuals, for which a convergence analysis is also given by using the martingale convergence theorem [3]. Nevertheless, none of these has been formally formulated in a standard online learning setting [4].

But in recent times there appeared several works [4–7] on developing online ARIMA algorithm for different applications.

Dataset Description

The first Dataset is Atmospheric CO_2 Data [8]. After dropping NaNs and splitting Dataset for train and test we get following:




Figure 1: Atmospheric CO_2 Dataset

Further we start analyzing of our Dataset:

- Estimation spectral density using Welch's method (figure 2):

The Welch method is based on Bartlett's method and differs in two ways: the signal is split up into overlapping segments: the original data segment is split up into L data segments of length M , overlapping by D points. If $D = M/2$, the overlap is said to be 50%. The overlapping segments are then windowed: After the data is split up into overlapping segments, the individual L data segments have a window applied to them (in the time domain). The windowing of the segments is what makes the Welch method a "modified" periodogram.

After doing the above, the periodogram is calculated by computing the discrete Fourier transform, and then computing the squared magnitude of the result. The individual periodograms are then averaged, which reduces the variance of the individual power measurements.

- Selection of the period for seasonal differences to delete seasonal component using SARIMA.
- Providing Augmented Dickey–Fuller (ADF) test:

The Dickey–Fuller test tests the null hypothesis that a unit root is present in an autoregressive model:

$$\Delta y_t = (\rho - 1)y_{t-1} + u_t = \delta y_{t-1} + u_t,$$

testing for a unit root is equivalent to testing $\delta = 0$. The testing procedure for the ADF test is the same as for the Dickey–Fuller test but it is applied to the model: $\Delta y_t = \alpha + \beta t + \gamma y_{t-1} + \delta_1 \Delta y_{t-1} + \dots + \delta_{p-1} \Delta y_{t-p+1} + \varepsilon_t$.

- Selection of the period for first differences to delete trend component.
- Providing Augmented Dickey–Fuller (ADF) test again to approve non-stationarity of the recieved Dataset.
- Searching for the hyperparameters (p,d,q), (P,D,Q,t) of the model SARIMAX by fitting a number of models and selecting the best depending on the AIC.

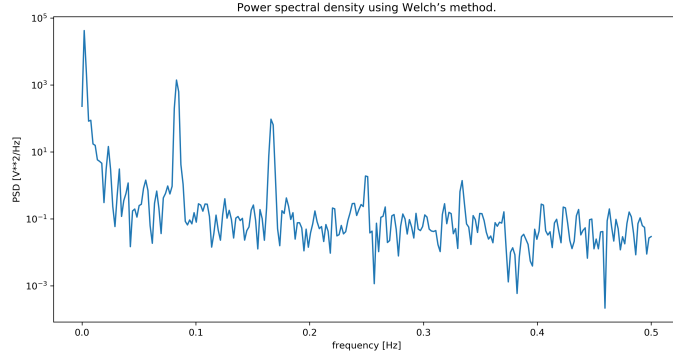


Figure 2: Spectral density estimation using Welch's method

The second Dataset is International airline passengers: monthly totals in thousands [9]. After dropping NaNs and splitting Dataset for train and test second dataset looks like at the figure 3. Then we provide all the procedures described earlier for the first Dataset.




Figure 3: International airline passengers: monthly totals in thousands

ML Methods and Experiments

Time series data are usually not realizations of a stationary process. For example, some of them may contain deterministic trends. An effective way to handle such strong serial correlations is to consider the differential method. For example, one can compute the first order differences of X_t by $\nabla X_t = X_t - X_{t-1}$ and the second order differences of X_t by $\nabla^2 X_t = \nabla X_t - \nabla X_{t-1}$.

If the sequence of $\nabla^d X_t$ satisfies an ARMA(k, q), we say that the sequence of X_t satisfies the ARIMA(k, d, q) (short for AutoRegressive Integrated Moving Average):

$$\nabla^d X_t = \sum_{i=1}^q \beta_i \epsilon_{t-i} + \sum_{i=1}^k \alpha_i \nabla^d X_{t-i} + \epsilon_t,$$

which are parameterized by three terms k, d, q and weights vector $\alpha \in \mathbb{R}^k$ and $\beta \in \mathbb{R}^q$. Note that ARMA(k,q) is a special case of the ARIMA(k, d, q), where the differences order is zero. Forecasting with ARIMA(k,d,q) is a reversion of differential process. Suppose time series sequence X_t satisfies ARIMA(k, d, q), we can predict the d-th order differential of observation at time $t + 1$ as \tilde{X}_t :

$$\tilde{X}_t = \nabla^d \tilde{X}_t + \sum_{i=0}^{d-1} \nabla^i X_{t-1}.$$

Consider an online ARIMA iteration at time t, the learner makes a prediction \tilde{X}_t and then the true X_t is disclosed to the learner. As a result, the learner suffers a loss $\ell_t(X_t, \tilde{X}_t)$. More formally, we can define the loss function as follows:

$$\begin{aligned} f_t(\alpha, \beta) &= \ell_t(X_t, \tilde{X}_t(\alpha, \beta)) = \ell_t\left(X_t, \left(\nabla^d \tilde{X}_t + \sum_{i=0}^{d-1} \nabla^i X_{t-1}\right)\right) \\ &= \ell_t\left(X_t, \left(\sum_{i=1}^q \beta_i \epsilon_{t-i} + \sum_{i=1}^k \alpha_i \nabla^d X_{t-i} + \sum_{i=0}^{d-1} \nabla^i X_{t-1}\right)\right) \end{aligned}$$

The goal of online ARIMA learning is to minimize the sum of losses over some number of rounds T . More formally, we can define the regret of the learner after T rounds as:

$$R_T = \sum_{t=1}^T \ell_t(X_t, \tilde{X}_t) - \min_{\alpha, \beta} \sum_{t=1}^T \ell_t(X_t, \tilde{X}_t(\alpha, \beta))$$

In our work we put into practice four different methods of online ARIMA implementa-

tion. The first three methods use built-in SARIMAX functions for fitting and forecasting from Statsmodel library for python 3.*. The fourth one

Method 1. Complete refitting

This method produces complete refitting at each step t of our model on the data of steps $[1, t]$.

Method 2. Rolling-window refitting

This method produces rolling-window refitting at each step t on the data of steps $[t - h, t]$.

Method 3. Warm start

This method uses built-in functions of statsmodel.SARIMAX by using params of the previously fitted model to speed up convergence of online fitting of the model at each step t .

Method 4. Gradient descent. Real online

We now implement a more general online convex optimization solver, Online Gradient Descent, which is applicable to any convex loss functions. Figure 4 presents the proposed ARIMA-OGD algorithm for optimizing the coefficient vector using the OGD algorithm. The projection $\Pi_{\mathcal{K}}(y)$ refers to the Euclidean projection onto \mathcal{K} , i.e., $\Pi_{\mathcal{K}}(y) = \arg \min_{x \in \mathcal{K}} \|y - x\|_2$.

Input: parameter k, d, q ; learning rate η .
Set $m = \log_{\lambda_{max}}((TLM_{max}q)^{-1})$.
for $t = 1$ **to** $T - 1$ **do**
 predict $\tilde{X}_t(\gamma^t) = \sum_{i=1}^{k+m} \gamma_i \nabla^d X_{t-i} + \sum_{i=0}^{d-1} \nabla^i X_{t-1}$;
 receive X_t and incur loss $\ell_t^m(\gamma^t)$;
 Let $\nabla_t = \nabla \ell_t^m(\gamma^t)$;
 Set $\gamma^{t+1} \leftarrow \Pi_{\mathcal{K}}(\gamma^t - \frac{1}{\eta} \nabla_t)$;
end for

Figure 4: ARIMA-OGD algorithm

All of the above in Method 4 section was an algorithm from [4]. We improved this algorithm by adding opportunity of using seasonal differences to delete seasonal component from the data.

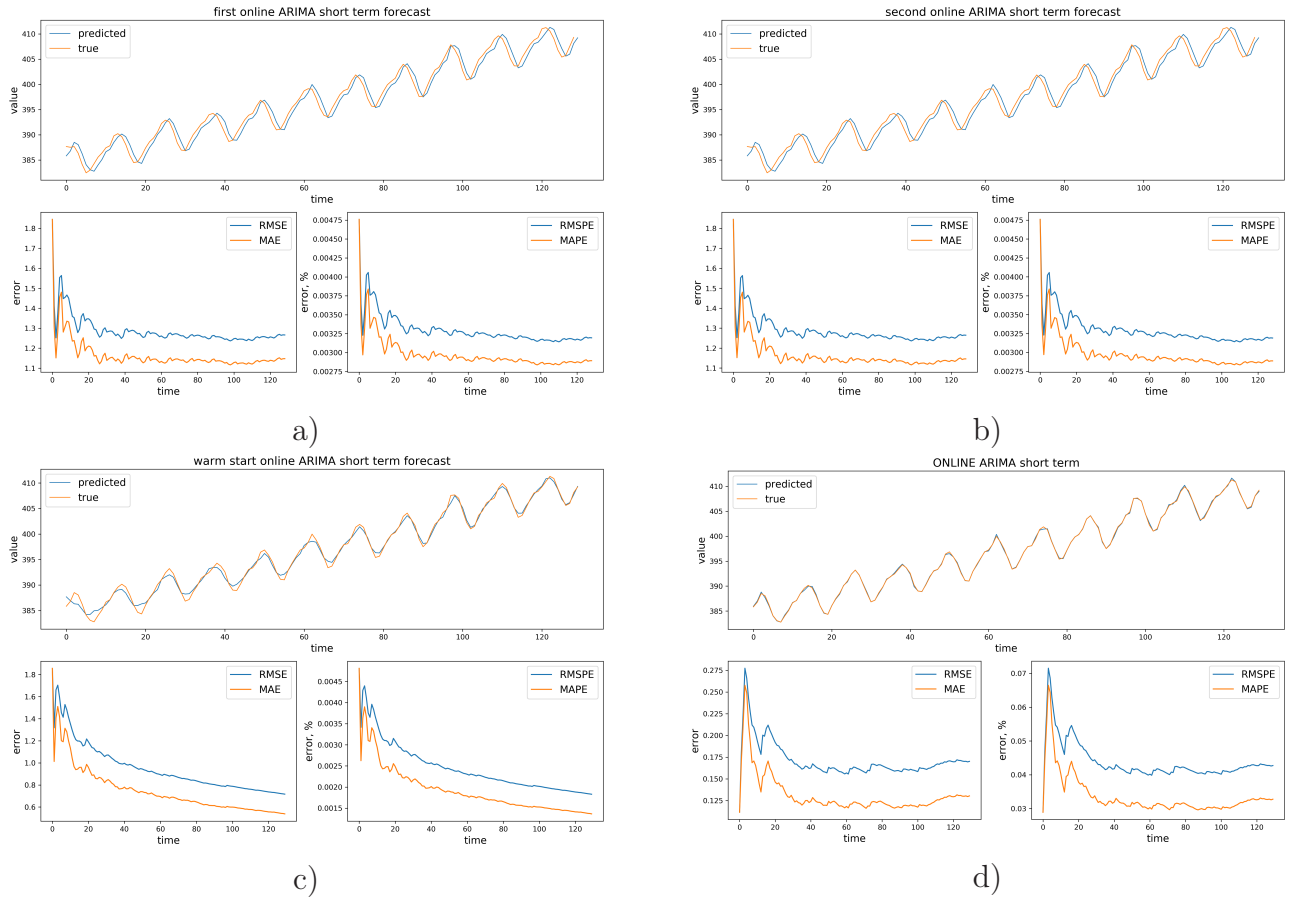


Figure 5: Example of short term (one point forward) forecast and errors: a) Complete refitting algorithm, b) Rolling-window refitting algorithm, c) Warm start algorithm, d) Gradient descent. Real online algorithm.

Conclusion

We empirically compared our four online SARIMA algorithms, in which the promising results on two different datasets validate that all of the algorithms are effective and promising for time series prediction. Though the fourth method has bigger MAPE and RMSPE, it has lower MAE and RMSE and visually better results. Our algorithms didn't show significant difference in forecasting accuracy but computational time in warm start case was on orders less (0.006 sec average for warm start algorithm vs 1.2 sec average for three others), and we can say that built-in function warm start is the best way to implement online time series forecasting for datasets with high frequency or, in other words, computationally complex problem.

Further research of this topic will include practical task analysis and improvements of implemented algorithms attempts for online: time reduction without reduced forecasting accuracy.

Bibliography

- [1] Huitian Lu, W.J. Kolarik, and S.S. Lu. Real-time performance reliability prediction. *IEEE Transactions on Reliability*, 50(4):353–357, 2001.
- [2] Ruey S. Tsay and George C. Tiao. Consistent estimates of autoregressive parameters and extended sample autocorrelation function for stationary and nonstationary ARMA models. *Journal of the American Statistical Association*, 79(385):84–96, mar 1984.
- [3] Feng Ding, Yang Shi, and Tongwen Chen. Performance analysis of estimation algorithms of nonstationary ARMA processes. *IEEE Transactions on Signal Processing*, 54(3):1041–1053, mar 2006.
- [4] P. Zhao C. Liu, S. C. H. Hoi and J. Sun. Online arima algorithms for time series prediction. In *Proceedings of AAAI Conference on Artificial Intelligence*, ser. AAAI’16, page pp. 1867–1873. AAAI Press, 2016.
- [5] Qiuwen Chen, Tiesheng Guan, Liu Yun, Ruonan Li, and Friedrich Recknagel. Online forecasting chlorophyll a concentrations by an auto-regressive integrated moving average model: Feasibilities and potentials. *Harmful Algae*, 43:58–65, mar 2015.
- [6] Johann Leithon, Teng Joon Lim, and Sumei Sun. Renewable energy management in cellular networks: An online strategy based on ARIMA forecasting and a markov chain model. In *2016 IEEE Wireless Communications and Networking Conference*. IEEE, apr 2016.
- [7] Florian Schmidt, Florian Suri-Payer, Anton Gulenko, Marcel Wallschlager, Alexander Acker, and Odej Kao. Unsupervised anomaly event detection for cloud monitoring using on-line arima. In *2018 IEEE/ACM International Conference on Utility and Cloud Computing Companion (UCC Companion)*. IEEE, dec 2018.
- [8] Atmospheric data: http://scrippsco2.ucsd.edu/data/atmospheric_co2/primary_monthly_record.
- [9] International airline passengers: monthly totals in thousands: <https://www.kaggle.com/andreazzini/international-airline-passengers/version/1>.